

EN3160 – Image Processing and Machine Vision

Assignment 2: Fitting and Alignment

Perera M.S.S - 220470K

GitHub Repository: <https://github.com/ShehanPer/EN3160-Assignment-02-220470K.git>

1 Question 1: Blob Detection

In this question, I used Laplacian of Gaussian (LoG) filters and scale-space extrema detection to identify blob regions in a sunflower field image.

Listing 1: Blob detection using LoG

```
1 sigma_values = np.linspace(1.0, 12.0, 12)
2 scale_space = []
3 for sigma in sigma_values:
4     hw = int(3 * sigma)
5     X, Y = np.meshgrid(np.arange(-hw, hw + 1), np.arange(-hw, hw + 1))
6     LoG = ((X**2 + Y**2) / (2 * sigma**2) - 1) * np.exp(-(X**2 + Y**2) /
7         (2 * sigma**2))
8     LoG = (LoG / (np.pi * sigma**4)) * sigma**2 # normalized LoG filter
9     response = cv.filter2D(gray, cv.CV_32F, LoG)
10    response = np.square(response)
11    scale_space.append(response) # store for different sigma
12 scale_space = np.stack(scale_space, axis=-1)
13 blobs = []
14 threshold = 0.1
15 for s in range(1, len(sigma_values)-1):
16     response = scale_space[:, :, s]
17     for i in range(1, response.shape[0]-1):
18         for j in range(1, response.shape[1]-1):
19             local_cube = scale_space[i-1:i+2, j-1:j+2, s-1:s+2]
20             if response[i, j] == np.max(local_cube) and response[i, j] >
21                 threshold:#local maxima detection
22                 radius = np.sqrt(2) * sigma_values[s]
23                 blobs.append((j, i, radius))
```

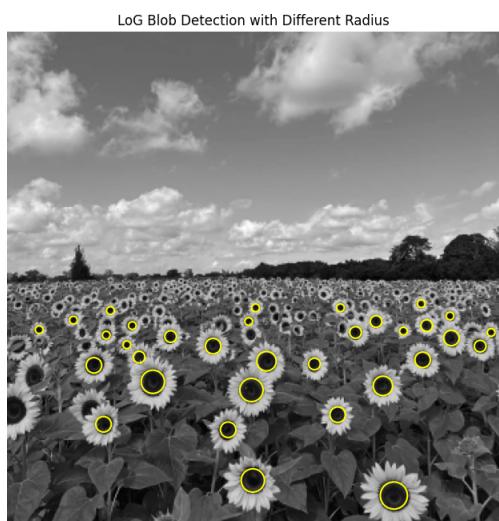


Figure 1 shows the output from the blob detection. The centers of the flowers show a near-perfect match with the blob filter, producing high responses in those areas with $\text{threshold} > 0.1$. It can also be seen how different sizes of Gaussians respond to sunflowers of different radius. By scanning through σ values from 1.0 to 12.0, sunflowers of various sizes were detected. The relationship between radius and σ is given by $\sigma = \text{radius}/\sqrt{2}$. The 3 largest detected blobs had radius of approximately 9.90, 8.49, and 7.07 pixels, corresponding to $\sigma \approx 7.00, 6.00$, and 5.00, respectively.

Figure 1: Detected blobs in the sunflower image with threshold = 0.1.

2 Question 2: RANSAC Line and Circle Fitting

RANSAC was used to robustly fit a line and a circle to a noisy dataset. The estimated line parameters were $a = -0.7258$, $b = -0.6879$, and $d = -1.6124$ with 39/100 inliers, closely matching the ground truth line ($m = -1$, $b = 2$). For the circle, the estimated center was $(1.8281, 3.3601)$ with radius $r = 10.1561$ and 37/61 inliers, which is close to the ground truth circle $(2, 3)$ with $r = 9.5037$. These results demonstrate that RANSAC effectively rejects outliers and accurately recovers the underlying geometric models.

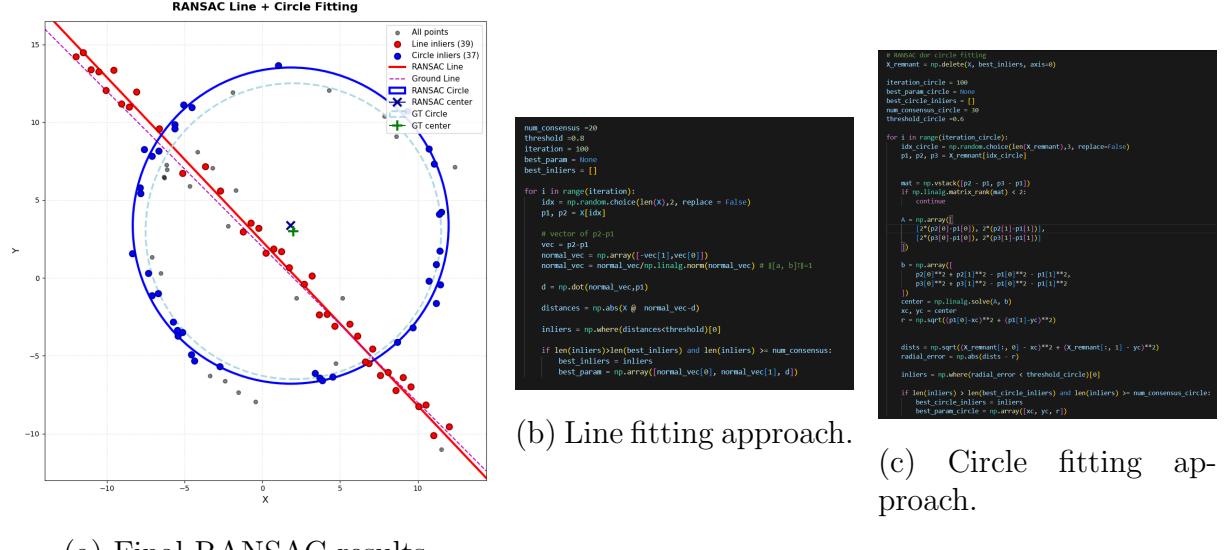


Figure 2: RANSAC-based model fitting results and corresponding implementation steps.

What Happens If the Circle is Fitted First?

If I try to fit the circle first, most of the line points become outliers, and the model struggles to find a proper circle. It can give wrong parameters or fail to detect the circle at all. Fitting the line first removes those points, so the circle fitting works much better afterward.

3 Question 3: Homography Transformation

For the homography transformation, two images were used, and points were selected via mouse clicks. Figures 3 (a) and (b) show the input and resulting overlay.

Listing 2: Computing and applying homography

```

1 main_img_points = np.array(img_points, dtype=np.float32)    # convert
   clicked points to array
2 flag_points = np.array([[0, 0], [flag_img.shape[1], 0], [flag_img.shape
   [1], flag_img.shape[0]], [0, flag_img.shape[0]]], dtype=np.float32)
3 H, _ = cv.findHomography(flag_points, main_img_points)    # compute
   homography matrix
4 warped_flag = cv.warpPerspective(flag_img, H, (main_img.shape[1],
   main_img.shape[0]))    # warp flag image to fit target region
5 mask = np.zeros_like(main_img, dtype=np.uint8)
6 cv.fillConvexPoly(mask, main_img_points.astype(int), (255, 255, 255))
7 flag_region = cv.bitwise_and(warped_flag, mask)
8 background = cv.bitwise_and(main_img, cv.bitwise_not(mask))
9 blend_img = cv.add(background, flag_region)    # combine warped flag with
   background

```



(a) Overlay result using homography.

(b) Second overlay result.

Figure 3: Comparison of overlay results using homography on two different images.

In this task, I selected pairs of images where one was superimposed onto a planar region of another. While choosing the images, I carefully considered their aspect ratios to ensure that the overlaid images did not appear stretched or distorted. Attention was also given to maintaining proper proportions and perspective alignment, resulting in a natural and visually consistent final composition.

4 Question 4: SIFT Feature Matching and Stitching

I computed SIFT features, matched them, estimated the homography within RANSAC, and stitched images. The following code was used to get the SIFT feature matching

Listing 3: SIFT feature matching

```

1 def get_sift_features(img1, img2, ratio_test=0.8):
2     img1_gray = cv2.cvtColor(img1, cv2.COLOR_BGR2GRAY)
3     img2_gray = cv2.cvtColor(img2, cv2.COLOR_BGR2GRAY)
4     sift = cv2.SIFT_create(nOctaveLayers=3, contrastThreshold=0.09,
5                           edgeThreshold=25, sigma=1) # Create SIFT detector
6     kp1, des1 = sift.detectAndCompute(img1_gray, None)
7     kp2, des2 = sift.detectAndCompute(img2_gray, None)
8     bf = cv2.BFMatcher() # Match descriptors using BFMatcher
9     matches_knn = bf.knnMatch(des1, des2, k=2)
10    best_matches = [m for m, n in matches_knn if m.distance < ratio_test
11                   * n.distance]
12    return best_matches, kp1, kp2

```



Figure 4: SIFT feature matching

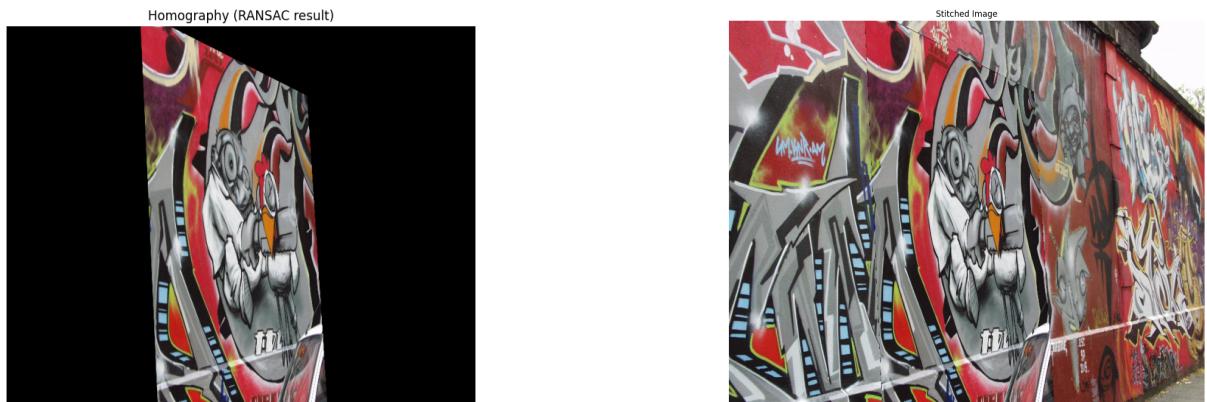
I used the RANSAC method to compute the best homography matrix, which robustly handles outliers in the matched points between images.

Listing 4: RANSAC-based homography

```

1 def find_Htransform(best_matches, keypoints1, keypoints5):
2     num_points, threshold, iterations = 4, 8, 500
3     best_homography, best_inlier_count, best_inliers = None, 0, None
4     all_test_points = []
5     all_dest_points = []
6     for match in best_matches:
7         all_test_points.append(np.array(keypoints1[match.queryIdx].pt))
8         all_dest_points.append(np.array(keypoints5[match.trainIdx].pt))
9     all_test_points = np.array(all_test_points)
10    all_dest_points = np.array(all_dest_points)
11    for i in range(iterations):
12        selected_idx = np.random.choice(len(best_matches), num_points,
13                                         replace=False)
14        selected_points = [best_matches[i] for i in selected_idx]
15        test_points = []
16        dest_points = []
17        for point in selected_points:
18            test_points.append(np.array(keypoints1[point.queryIdx].pt))
19            dest_points.append(np.array(keypoints5[point.trainIdx].pt))
20        test_points = np.array(test_points)
21        dest_points = np.array(dest_points)
22        h_transform = transform.estimate_transform("projective",
23                                                test_points, dest_points)
23        calc_points = h_transform(all_test_points)
24        errors = np.sqrt(np.sum(np.square(calc_points-all_dest_points),
25                               axis=1))
26        inliers = np.where(errors<threshold)[0]
27        if len(inliers)> best_inlier_count:
28            best_inlier_count = len(inliers)
29            best_inliers = inliers
30            best_homography = h_transform
31    return best_homography, best_inliers

```



(a) Homography transformation.

(b) Stitched image.

Figure 5: Homography Transformation and stitched image

The homography matrix computed using RANSAC is similar to the provided dataset matrix (H1to5p) :

$$H_{\text{computed}} = \begin{bmatrix} 0.6461 & 0.0322 & 227.36 \\ 0.2468 & 1.1410 & -24.32 \\ 0.00057 & -0.000086 & 1 \end{bmatrix}, \quad H_{\text{given}} = \begin{bmatrix} 0.6254 & 0.0578 & 222.01 \\ 0.2224 & 1.1652 & -25.61 \\ 0.00049 & -0.0000365 & 1 \end{bmatrix}.$$